

# What you write is not what SQL sees

IMPROVING QUERY PERFORMANCE IN POSTGRESQL



Amy McCarty  
Instructor

# Algebraic order of operations

- Lexical (as written)
- Logical (as executed)

PEMDAS	BODMAS
Parenthesis	Brackets
Exponents	Order
Multiplication /Division	Division /Multiplication
Addition /Subtraction	Addition /Subtraction

# Applying the order of operations

Lexical:

$$x = 2 + (8 + 4) \div 2$$

$$x = 10 + 4) \div 2$$

$$x = 14 \div 2$$

$$x = 7$$

Logical:

$$x = 2 + (8 + 4) \div 2$$

$$x = 2 + \frac{12}{2}$$

$$x = 2 + 6$$

$$x = 8$$

# SQL logical order of operations

Order	Clause	Purpose
1	<b>FROM</b>	Provides directions to the table or tables if the query includes joins
2	<b>WHERE</b>	Filters or limits the records
3	<b>GROUP BY</b>	Places records into categories
4	<b>SUM(), COUNT(), etc</b>	Aggregates
5	<b>SELECT</b>	identifies columns to return

```
SELECT COUNT(*) FROM tableA WHERE col1 = 77
```

# Group by and aggregations

event_location	storm	elements	days
Russia	blizzard	water	1
Argentina	tornado	water	1
Argentina	tornado	wind	1
Australia	tornado	wind	1
Kuwait	haboob	wind	2
USA	haboob	wind	2

```
SELECT elements, storm, COUNT(*)  
FROM weather_events  
GROUP BY elements
```

No output -  
your code generated an error

column "storm" must appear in the  
GROUP BY clause or be used in an  
aggregate function

# Group by and aggregations order of operations

```
SELECT elements, storm, COUNT(*)  
FROM weather_events  
GROUP BY elements
```

order	SQL clause	available columns
1	FROM weather_events	all
3	GROUP BY elements	elements
4	COUNT	elements

# Group by matches the aggregations

```
SELECT elements, storm, COUNT(*)  
FROM weather_events  
GROUP BY elements, storm
```

# Group by matches the aggregations

```
SELECT elements, storm, COUNT(*)  
FROM weather_events  
GROUP BY elements, storm
```

## Results

elements	storm	count
water	blizzard	1
water	tornado	1
wind	tornado	2
wind	haboob	2



# SQL logical order of operations continued

Order	Clause	Purpose
...		
5	<b>SELECT</b>	identifies columns to return
6	<b>DISTINCT</b>	removes duplicates
7	<b>ORDER BY</b>	arranges results
8	<b>LIMIT</b>	removes rows

# Distinct and limit

row_no	location	storm	elements	days
1	Russia	blizzard	water	1
2	Argentina	tornado	water	1
3	Argentina	tornado	wind	1
4	Australia	tornado	wind	1
5	Kuwait	haboob	wind	2
6	USA	haboob	wind	2

```
SELECT DISTINCT storm, elements
FROM weather_events
ORDER BY storm LIMIT 3
```

row_no	location	storm	elements	days
1	Russia	blizzard	water	1
2	Argentina	tornado	water	1
3	Argentina	tornado	wind	1
4	Australia	tornado	wind	1
5	Kuwait	haboob	wind	2
6	USA	haboob	wind	2

```
SELECT DISTINCT storm, elements
FROM weather_events
ORDER BY storm LIMIT 3
```

row_no	location	storm	elements	days
1	Russia	blizzard	water	1
2	Argentina	tornado	water	1
3	Argentina	tornado	wind	1
4	Australia	tornado	wind	1
5	Kuwait	haboob	wind	2
6	USA	haboob	wind	2

order	SQL clause	available rows
1	FROM weather_events	all

```

SELECT DISTINCT storm, elements
FROM weather_events
ORDER BY storm LIMIT 3

```

row_no	storm	elements
1	blizzard	water
2	tornado	water
3	tornado	wind
5	haboob	wind

order	SQL clause	available rows
1	FROM weather_events	all
5	SELECT storm, elements	all

```
SELECT DISTINCT storm, elements
FROM weather_events
ORDER BY storm LIMIT 3
```

row_no	storm	elements
1	blizzard	water
2	tornado	water
3	tornado	wind
5	haboob	wind

order	SQL clause	available rows
1	FROM weather_events	all
5	SELECT storm, elements	all
6	DISTINCT	1, 2, 3, 5

```
SELECT DISTINCT storm, elements
FROM weather_events
ORDER BY storm LIMIT 3
```

row_no	storm	elements
1	blizzard	water
5	haboob	wind
2	tornado	water
3	tornado	wind

order	SQL clause	available rows
1	FROM weather_events	all
5	SELECT storm, elements	all
6	DISTINCT	1, 2, 3, 5
7	ORDER BY storm	1, 5, 2, 3

```
SELECT DISTINCT storm, elements
FROM weather_events
ORDER BY storm LIMIT 3
```

row_no	storm	elements
1	blizzard	water
5	haboob	wind
2	tornado	water

order	SQL clause	available rows
1	FROM weather_events	all
5	SELECT storm, elements	all
6	DISTINCT	1, 2, 3, 5
7	ORDER BY storm	1, 5, 2, 3
8	LIMIT 3	1, 5, 2

```
SELECT DISTINCT storm, elements
FROM weather_events
ORDER BY storm LIMIT 3
```



Order	Clause	Purpose	Limits
1	<b>FROM</b>	provides directions to the table(s)	
2	<b>WHERE</b>	filters or limits the records	# rows
3	<b>GROUP BY</b>	places records into categories	# columns
4	<b>SUM, COUNT, etc</b>	aggregates	# rows
5	<b>SELECT</b>	identifies columns to return	# columns
6	<b>DISTINCT</b>	removes duplicates	# rows
7	<b>ORDER BY</b>	arranges results	
8	<b>LIMIT</b>	filters records	# rows

# Let's practice!

IMPROVING QUERY PERFORMANCE IN POSTGRESQL

# Filtering in the **WHERE** clause

IMPROVING QUERY PERFORMANCE IN POSTGRESQL



Amy McCarty  
Instructor

# Limit the data

Order	Clause	Purpose
1	FROM	provides directions to the table(s)
2	<b>WHERE</b>	<b>filters or limits the records</b>

- Occurs early
- Fewer records

# EXPLAIN

**EXPLAIN**

**SELECT** \* **FROM** phones

```
Seq Scan on phones (cost = 0.00..22.7  
                    ,rows=1270  
                    ,width=36)
```

- Number of execution steps

Query planner



# EXPLAIN with WHERE

**EXPLAIN**

**SELECT** \* **FROM** phones

**WHERE** phone\_code = 235

```
Seq Scan on phones (cost = 0.00..25.8  
                    ,rows=6,width=636)  
  Filter: (phone_code=235)
```

Query planner



# Good - Filtering for similar values with LIKE OR

country	phone_code	reliability
Chad	235	medium
China	86	high
Costa Rica	506	high
India	91	medium
Indonesia	62	medium
Iraq	964	low

**EXPLAIN**

```
SELECT * FROM phones
WHERE country LIKE 'Ch%'
       OR country LIKE 'In%'
```

```
Seq Scan on phones (cost = 0.00..29.05
                    ,rows=13,width=36)
```

```
Filter: ((country~~'Ch%'::text)
        OR(country~~'In%'::text))
```

# Better - Filtering for similar values with LIKE ANY

country	phone_code	reliability
Chad	235	medium
China	86	high
Costa Rica	506	high
India	91	medium
Indonesia	62	medium
Iraq	964	low

# EXPLAIN

**SELECT** \* **FROM** phones

**WHERE** country

**LIKE ANY**(ARRAY[ 'Ch%', 'In%' ])

```
Seq Scan on phones (cost = 0.00..25.88  
                                , rows=13, width=36)
```

```
Filter: ((country~~ANY(' {Ch%,In%}'
                        ::text[ ]))
```



# Good - Filtering for exact values with OR

country	phone_code	reliability
Chad	235	medium
China	86	high
Costa Rica	506	high
India	91	medium
Indonesia	62	medium
Iraq	964	low

**EXPLAIN**

```
SELECT * FROM phones
WHERE country = 'Chad'
       OR country = 'China'
```

```
Seq Scan on phones (cost = 0.00..29.05
                    ,rows=13,width=36)
   Filter: ((country='Chad'::text)
           OR(country='China'::text))
```

# Better - Filtering for exact values with IN

country	phone_code	reliability
Chad	235	medium
China	86	high
Costa Rica	506	high
India	91	medium
Indonesia	62	medium
Iraq	964	low

**EXPLAIN**

**SELECT** \* **FROM** phones

**WHERE** country **IN** ( 'Chad' , 'China' )

```
Seq Scan on phones (cost = 0.00..25.88
                        , rows=13, width=36)
   Filter: ((country=ANY( ' {Chad,China}'
                          ::text[]))
```

# Best - Filtering for numbers

country	phone_code	reliability
Chad	235	medium
China	86	high
Costa Rica	506	high
India	91	medium
Indonesia	62	medium
Iraq	964	low

```
EXPLAIN
SELECT *
FROM phones
WHERE phone_code IN (235, 86)
```

```
Seq Scan on phones (cost = 0.00..25.88
                    ,rows=13,width=36)
   Filter: (phone_code=ANY(' {235,86}'
                           ::integer[]))
```

# Summarizing the best WHERE filters

## Numeric advantages

- Shorter length
- Smaller storage
- Speeds performance

# Summarizing the best WHERE filters

## Numeric advantages

- Shorter length
- Smaller storage
- Speeds performance

Good	Better
Text	Numeric
OR	IN, ARRAY

# Let's practice!

IMPROVING QUERY PERFORMANCE IN POSTGRESQL

# Filtering while joining

IMPROVING QUERY PERFORMANCE IN POSTGRESQL



Amy McCarty  
Instructor

# Joins revisited

## Joins

- Link tables using at least 1 common field

```
SELECT *  
FROM TABLE_A AS A  
INNER JOIN TABLE_B AS B  
ON A.NAME = B.NAME
```

## Joins to combine data

- From multiple tables
- Inner and outer

## Joins to filter data

- Inner join
- Outer join with non-linking join condition



# Patient and appointments data

Appointments

visit_id	reason	patient_id
01	checkup	999
02	infection	888

Patients

patient_id	name	sex
999	Lotte Smith	F
888	Zhang Wei	M
777	Amelia Hernandez	F

# Inner joins to filter


```
SELECT *  
FROM appointments a  
INNER JOIN patients p  
ON a.patient_id = p.patient_id
```

Appointments

visit_id	reason	patient_id
01	checkup	999
02	infection	888

Patients

patient_id	name	sex
999	Lotte Smith	F
888	Zhang Wei	M
<del>777</del>	<del>Amelia Hernandez</del>	<del>F</del>



# Inner joins to filter

```
SELECT *  
FROM appointments a  
INNER JOIN patients p  
ON a.patient_id = p.patient_id
```

## Output

visit_id	reason	patient_id	name	sex
01	checkup	999	Lotte Smith	F
02	infection	888	Zhang Wei	M

# Outer joins to filter

```
SELECT *  
FROM appointments a  
LEFT JOIN patients p  
  ON a.patient_id = p.patient_id  
  AND p.sex = 'M'
```

Appointments

visit_id	reason	patient_id
01	checkup	999
02	infection	888

Patients

patient_id	name	sex
999	Lotte Smith	F
888	Zhang Wei	M
777	Amelia Hernandez	F



# Outer joins to filter

```
SELECT *  
FROM appointments a  
LEFT JOIN patients p  
  ON a.patient_id = p.patient_id  
  AND p.sex = 'M'
```

## Output

visit_id	reason	patient_id	name	sex
01	checkup	999		
02	infection	888	Zhang Wei	M

# Filter pitfalls

```
SELECT *  
FROM appointments a  
LEFT JOIN patients p  
  ON a.patient_id = p.patient_id  
WHERE p.sex = 'M'
```

## 1) FROM

Appointments

visit_id	reason	patient_id
01	checkup	999
02	infection	888

Patients

patient_id	name	sex
999	Lotte Smith	F
888	Zhang Wei	M
777	Amelia Hernandez	F

## 2) WHERE

Output

visit_id	reason	patient_id	name	sex
<del>01</del>	<del>checkup</del>	<del>999</del>	<del>Lotte Smith</del>	<del>F</del>
02	infection	888	Zhang Wei	M

# Filter pitfalls

```
SELECT *  
FROM appointments a  
LEFT JOIN patients p  
  ON a.patient_id = p.patient_id  
WHERE p.sex = 'M'
```

## Output

visit_id	reason	patient_id	name	sex
02	infection	888	Zhang Wei	M

# Filter pitfalls

```
SELECT *  
FROM appointments a  
LEFT JOIN patients p  
  ON a.patient_id = p.patient_id  
WHERE p.sex = 'M'
```

## Output

visit_id	reason	patient_id	name	sex
02	infection	888	Zhang Wei	M

```
SELECT *  
FROM appointments a  
LEFT JOIN patients p  
  ON a.patient_id = p.patient_id  
AND p.sex = 'M'
```

visit_id	reason	patient_id	name	sex
01	checkup	999		
02	infection	888	Zhang Wei	M



# Filter pitfalls improved

```
SELECT *  
FROM appointments a  
INNER JOIN patients p  
    ON a.patient_id = p.patient_id  
WHERE p.sex = 'M'
```

```
SELECT *  
FROM appointments a  
INNER JOIN patients p  
    ON a.patient_id = p.patient_id  
AND p.sex = 'M'
```

visit_id	reason	patient_id	name	sex	
02	infection	888	Zhang Wei	M	

# Let's practice!

IMPROVING QUERY PERFORMANCE IN POSTGRESQL

# Aggregating with different data granularities

IMPROVING QUERY PERFORMANCE IN POSTGRESQL



Amy McCarty  
Instructor

# Data granularity - level of detail

- Makes a row unique
- 1+ columns

Video\_Games

id	game	first_yr	
012	Grand Theft Auto	1997	
234	Legend of_Zelda	1986	

Game\_Platforms

game_id	platform	year	
234	FCDS	1986	
234	GameCube	2003	
234	Wii	2006	


# Joining different granularities

Video\_Games

id	game	first_yr
012	Grand Theft Auto	1997
234	Legend of_Zelda	1986

Game\_Platforms

game_id	platform	year
234	FCDS	1986
234	GameCube	2003
234	Wii	2006



id	game	first_yr	no_platforms
234	Legend_of_Zelda	1986	3
234	Legend_of_Zelda	1986	3
234	Legend_of_Zelda	1986	3

```
SELECT g.id, g.game, g.first_yr
      , COUNT(platform) AS no_platforms
FROM video_games g
INNER JOIN game_platforms p
      ON g.id = p.game_id
GROUP BY g.id, g.game, g.first_yr
```

# Changing the table granularity

## Game\_Platforms

game_id	platform	year
234	FCDS	1986
234	GameCube	2003
234	Wii	2006

```
SELECT game_id
       , COUNT(platform) as no_platforms
FROM game_platforms
GROUP BY game_id
```

# Changing the table granularity

## Game\_Platforms

game_id	platform	year
234	FCDS	1986
234	GameCube	2003
234	Wii	2006

```
SELECT game_id
, COUNT(platform) as no_platforms
FROM game_platforms
GROUP BY game_id
```

## Output

game_id	no_platforms
234	3

# CTEs revisited

- Standalone query with temporary results set
- With statement structure

```
WITH cte AS  
  ( SELECT * FROM tableA )  
  
SELECT * FROM cte
```

- Allows aggregation before joining



# CTEs to the granularity rescue

```
WITH platforms_cte AS  
(  
  SELECT game_id  
        , COUNT(platforms) as no_platforms  
  FROM game_platforms  
  GROUP BY game_id  
)
```

```
SELECT g.id, g.game , cte.no_platforms  
FROM video_games g  
INNER JOIN platforms_cte cte  
  ON g.id = cte.game_id
```

# CTEs to the granularity rescue

```
WITH platforms_cte AS
(
  SELECT game_id
        , COUNT(platforms) as no_platforms
  FROM game_platforms
  GROUP BY game_id
)
```

```
SELECT g.id, g.game , cte.no_platforms
FROM video_games g
INNER JOIN platforms_cte cte
  ON g.id = cte.game_id
```

Output

id	game	no_platforms
234	Legend of Zelda	3

# Matching data granularity when joining

- No repeats or duplicates
- Minimum needed results
- No double counting

# Let's practice!

IMPROVING QUERY PERFORMANCE IN POSTGRESQL